

INTERPROCESS COMMUNICATION

**From Chapter 4 of Distributed Systems
Concepts and Design, 4th Edition,**

By G. Coulouris, J. Dollimore and T. Kindberg

**Published by Addison Wesley/Pearson
Education June 2005**

Topics

- INTRODUCTION
- The API for the INTERNET PROTOCOLS
- EXTERNAL DATA REPRESENTATION
- CLIENT-SERVER COMMUNICATION

4.1 Introduction

- A process can be: **Independent** process or **Co-operating process**.
- An **independent process** is **not affected** by the execution of other processes while a **co-operating process** can be **affected** by other executing processes.
- Processes running **independently**, will execute **very efficiently** but in practical, there are many situations when **co-operative nature** can be utilized for **increasing computational speed, convenience and modularity**.
- Inter process communication (IPC)** is a mechanism which allows processes to **communicate** each other and **synchronize** their actions.

Processes can communicate with each other using these two ways : **Shared Memory** or **Message passing**

4.1 Introduction

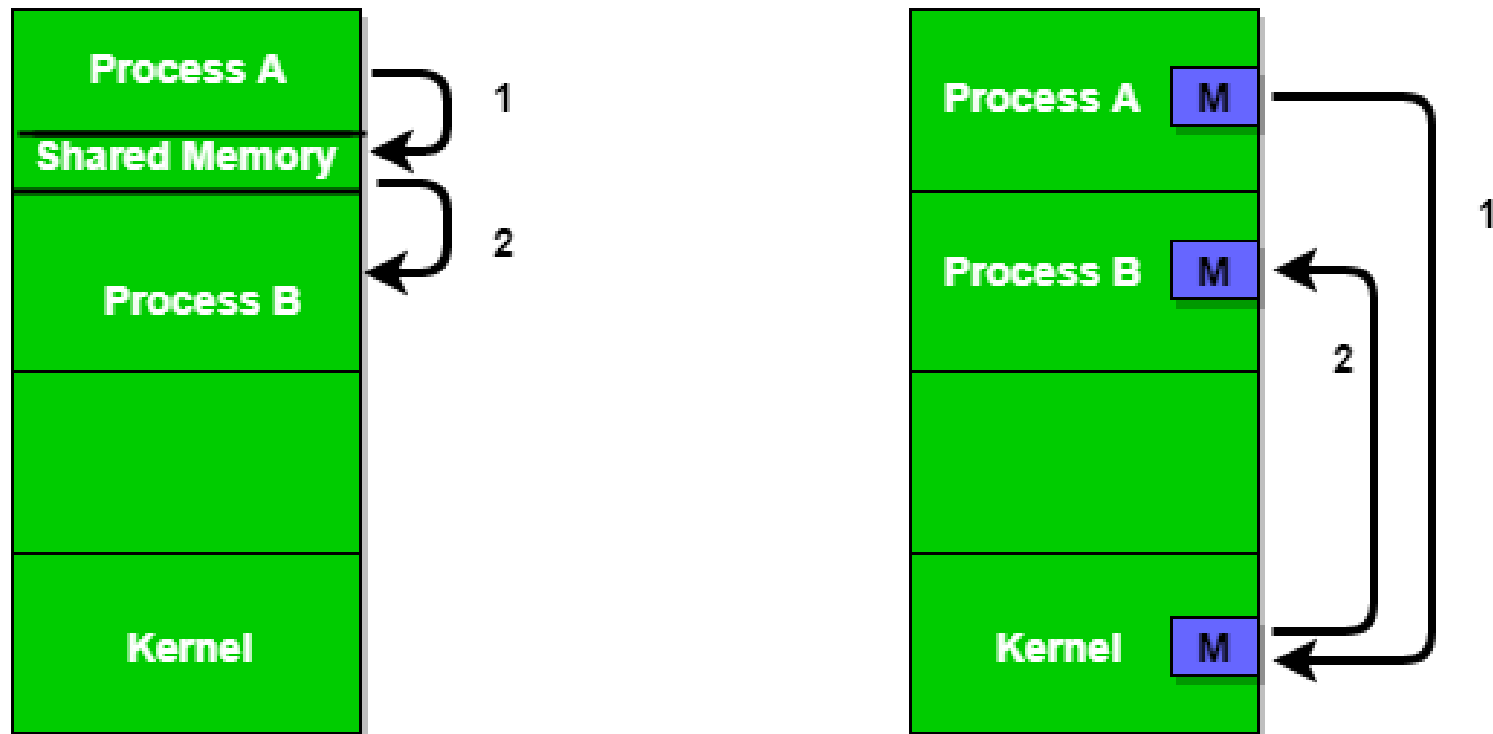


Figure 1 - Shared Memory and Message Passing

4.1 Introduction

- The java API for interprocess communication in the internet provides both **datagram** and **stream** communication.
- The communication patterns that are most commonly used in distributed programs:

- **Client-Server communication**

- ❖ The request and reply messages provide the basis for remote method invocation (RMI) or remote procedure call (RPC).

- **Group communication**

- ❖ The same message is sent to several processes
- ❖ **Group multicast** communication in which one process in a group transmits the same message to all members of the group

Introduction

- This chapter is concerned with middleware.

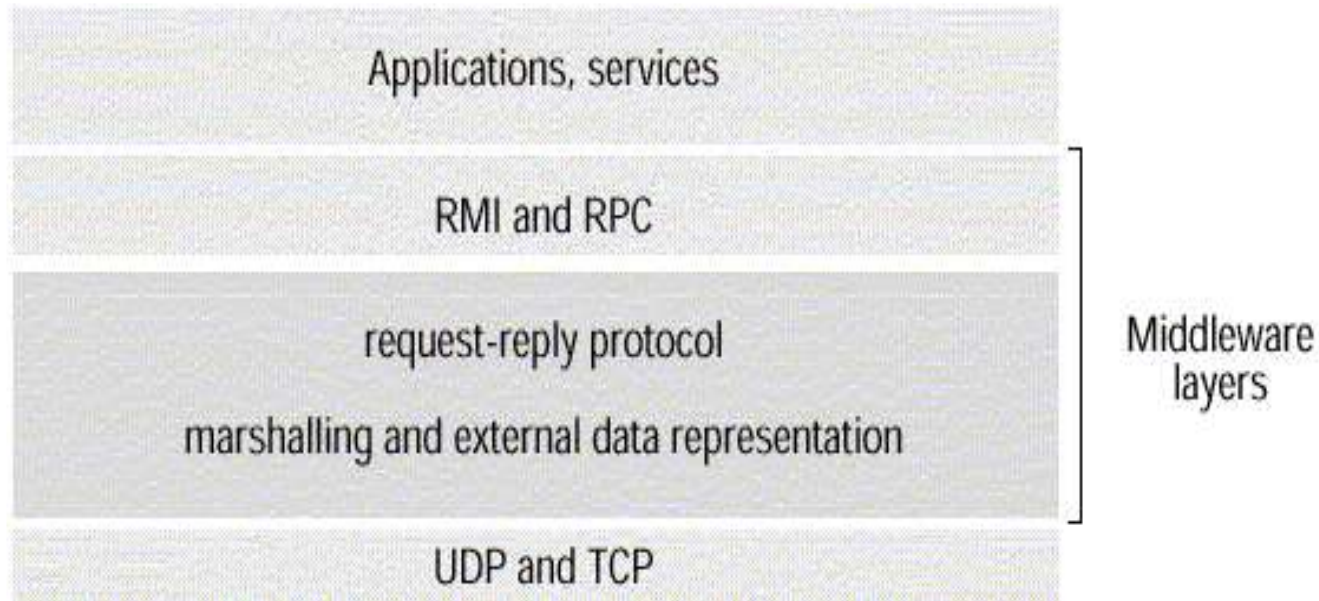


Figure 1. Middleware layers

Introduction

- Remote Method Invocation (RMI)

- It allows an object to invoke a method in an object in a remote process.

- ❖ E.g. CORBA and Java RMI

- Remote Procedure Call (RPC)

- It allows a client to call a procedure in a remote server.

Introduction

- The application program interface (API) to UDP provides a **message passing** abstraction.
 - Message passing is the simplest form of interprocess communication.
 - API enables a sending process to transmit a single message to a receiving process.
 - The independent packets containing these messages are called **datagrams**.
 - In the Java and UNIX APIs, the sender specifies the destination using a **socket**.

Introduction

- **Socket** is an indirect reference to a particular port used by the destination process at a destination computer.
- The application program interface (API) to TCP provides the abstraction of a two-way stream between pairs of processes.
- The information communicated consists of a stream of data items with no message boundaries.

4.2 The API for the Internet Protocols

- The CHARACTERISTICS of INTERPROCESS COMMUNICATION
- SOCKET
- UDP DATAGRAM COMMUNICATION
- TCP STREAM COMMUNICATION

4.2.1 The Characteristics of Interprocess Communication

1. Synchronous and asynchronous communication

- In the synchronous form, both send and receive are **blocking** operations.
Eg. Continuous Chatting
- In the asynchronous form, the use of the **send** operation is **non-blocking** and the **receive** operation can have **blocking** and **non-blocking** variants.
Eg. WhatsApp

The Characteristics of Interprocess Communication

2. Message destinations

- A local port is a message destination within a computer, specified as an integer.
- A port has an exactly one receiver but can have many senders.

The Characteristics of Interprocess Communication

3. Reliability

- A reliable communication is defined in terms of **validity** and **integrity**.
- A **point-to-point message** service is described as reliable if messages are **guaranteed to be delivered** despite a reasonable number of packets being dropped or lost.
- For integrity, messages must arrive **uncorrupted** and **without duplication**.


4. Ordering: Some applications require that messages to be delivered in sender order.

4.2 Sockets

- A **socket** can be thought of as an endpoint in a two-way communication channel. Eg:telephone Call
- Internet IPC mechanism of Unix and other operating systems (BSD Unix, Solaris, Linux, Windows NT, Macintosh OS) Processes can **send** and **receive** messages **via a socket**.
- Sockets need to be bound to a **port number and an internet address** in order to send and receive messages.
- Each socket has a transport layer protocol (TCP or UDP).

Sockets

- Messages sent to some internet address and port number can only be received by a process using a socket that is bound to this address and port number.



IP Address + Port number = Socket

TCP/IP Ports And Sockets

Sockets

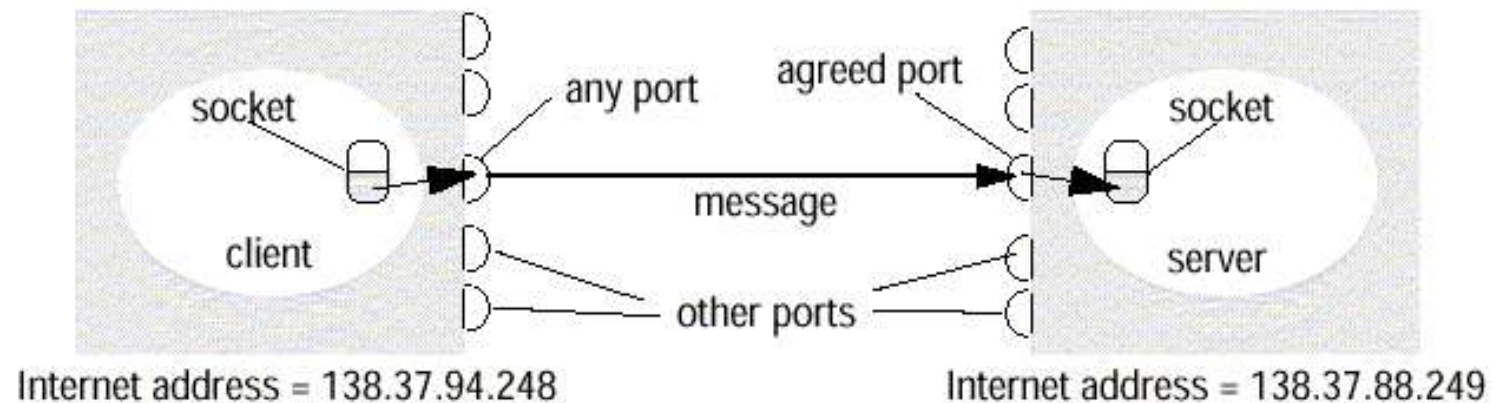


Figure 2. Sockets and ports

4.2.1 UDP Datagram Communication

- UDP datagram properties
 - No guarantee of order preservation
 - Message loss and duplications are possible
- Necessary steps
 - Creating a socket
 - Binding a socket to a port and local Internet address
 - ❖ A client binds to any free local port
 - ❖ A server binds to a server port
- Receive method
 - It returns Internet address and port of sender, plus message.

UDP Datagram Communication

- Issues related to datagram communications are:
 - Message size
 - ❖ IP allows for messages of up to 2^{16} bytes.
 - ❖ Most implementations restrict this to around 8k bytes.
 - ❖ Any application requiring messages larger than the maximum must fragment.
 - ❖ If arriving message is too big for array allocated to receive message content, truncation occurs.

UDP Datagram Communication

➤ Blocking

❖ Send: non-blocking

- upon arrival, message is placed in a queue for the socket that is bound to the destination port.

❖ Receive: blocking

- Pre-emption by timeout possible
- If process wishes to continue while waiting for packet, use separate thread

➤ Timeout

➤ Receive from any

UDP Datagram Communication

- UDP datagram's suffer from following failures:
 - Omission failure
 - Messages may be dropped occasionally
 - Ordering

Java API for UDP Datagrams

- The Java API provides datagram communication by two classes:

- Datagram Packet

- ❖ It provides a constructor to make an array of bytes comprising:
 - Message content
 - Length of message
 - Internet address
 - Local port number

array of bytes containing message | length of message| Internet address | port number|

- ❖ It provides another similar constructor for receiving a message.

Java API for UDP Datagrams

➤ Datagram Socket

- ❖ This class supports sockets for sending and receiving UDP datagram.
- ❖ It provides a constructor with port number as argument.

- ❖ Datagram Socket methods are:
 - `send` and `receive`
 - `setSoTimeout`
 - `connect`

4.2.2 TCP Stream Communication

- The API to the TCP protocol provides the abstraction of a stream of bytes to be written to or read from.
 - Characteristics of the stream abstraction:
 - ❖ Message sizes: Application dependent
 - ❖ Lost messages: Ack, Seq No.
 - ❖ Flow control: Window size
 - ❖ Message duplication & Ordering : Seq No
 - ❖ Message destination: One time with IP & Port no.

TCP Stream Communication

- Issues related to stream communication:
 - Matching of data items:
 - ❖ data interpretation error w.r.t. use of data stream
 - Blocking: Use of Queue destination Socket
 - Threads: To avoid delay in handling clients

TCP Stream Communication

■ Use of TCP

➤ Many services that run over TCP connections, with reserved port number are:

- ❖ HTTP (Hypertext Transfer Protocol)
- ❖ FTP (File Transfer Protocol)
- ❖ Telnet
- ❖ SMTP (Simple Mail Transfer Protocol)

TCP Stream Communication

■ Java API for TCP streams

➤ The Java interface to TCP streams is provided in the classes:

❖ ServerSocket

- It is used by a server to create a socket at server port to listen for connect requests from clients.

❖ Socket

- It is used by a pair of processes with a connection.
- The client uses a constructor to create a socket and connect it to the remote host and port of a server.
- It provides methods for accessing input and output streams associated with a socket.

4.3 External Data Representation

- The information stored in running programs is represented as data structures, whereas the information in messages consists of sequences of bytes.
- Irrespective of the form of communication used, the data structure must be converted to a sequence of bytes before transmission and rebuilt on arrival.

External Data Representation

- External Data Representation is an agreed standard for the representation of data structures and primitive values.
- Data representation problems are:
 - Using agreed external representation, two conversions necessary
 - Using sender's or receiver's format and convert at the other end

External Data Representation

- **Marshalling**
 - Marshalling is the process of taking a collection of data items and assembling them into a form suitable for transmission in a message.
- **Unmarshalling**
 - Unmarshalling is the process of disassembling a collection of data on arrival to produce an equivalent collection of data items at the destination.

External Data Representation

- Three approaches to external data representation and marshalling are:
 - CORBA
 - Java's object serialization

External Data Representation

- Marshalling and unmarshalling activities is usually performed automatically by middleware layer.
- Marshalling is likely error-prone if carried out by hand.

4.3.1 CORBA Common Data Representation (CDR)

- CORBA Common Data Representation (CDR)
 - CORBA CDR is the external data representation defined with CORBA 2.0.
 - It consists 15 primitive types:
 - Short (16 bit)
 - Long (32 bit)
 - Unsigned short
 - Unsigned long
 - Float(32 bit)
 - Double(64 bit)
 - Char
 - Boolean(TRUE,FALSE)
 - Octet(8 bit)
 - Any(can represent any basic or constructed type)
 - Composite type are shown in Figure 8.

CORBA Common Data Representation (CDR)

| <i>Type</i> | <i>Representation</i> |
|-------------------|--|
| <i>sequence</i> | length (unsigned long) followed by elements in order |
| <i>string</i> | length (unsigned long) followed by characters in order (can also can have wide characters) |
| <i>array</i> | array elements in order (no length specified because it is fixed) |
| <i>struct</i> | in the order of declaration of the components |
| <i>enumerated</i> | unsigned long (the values are specified by the order declared) |
| <i>union</i> | type tag followed by the selected member |

Figure 8. CORBA CDR for constructed types

CORBA Common Data Representation (CDR)

Figure 9 shows a message in CORBA CDR that contains the three fields of a **struct** whose respective types are **string**, **string**, and **unsigned long**.

example: struct with value {'Smith', 'London', 1934}

| <i>index in sequence of bytes</i> | | <i>notes on representation</i> |
|---------------------------------------|----------|------------------------------------|
| 0–3 | 5 | <i>length of string</i> |
| 4–7 | "Smit" | 'Smith' |
| 8–11 | "h____" | |
| 12–15 | 6 | <i>length of string</i> |
| 16–19 | "Lond" | 'London' |
| 20–23 | "on____" | |
| 24–27 | 1934 | <i>unsigned long</i> |

Figure 9. CORBA CDR message

4.3.2 Java object serialization

- In Java RMI, both object and primitive data values may be passed as arguments and results of method invocation.
- An object is an instance of a Java class.
 - Example, the Java class equivalent to the Person struct

```
Public class Person implements Serializable {  
    Private String name;  
    Private String place;  
    Private int year;  
    Public Person(String aName ,String aPlace, int aYear) {  
        name = aName;  
        place = aPlace;  
        year = aYear;  
    }  
    //followed by methods for accessing the instance variables  
}
```

Java object serialization

The serialized form is illustrated in **Figure 10**.

| <i>Serialized values</i> | | | | <i>Explanation</i> |
|--------------------------|-----------------------|--------------------------|---------------------------|--|
| Person | 8-byte version number | | h0 | class name, version number |
| 3 | int year | java.lang.String name | java.lang.String place | number, type and name of instance variables |
| 1934 | 5 Smith | 6 London | h1 | values of instance variables |

Figure 10. Indication of Java serialization form

Remote Object References

- Remote object references are needed when a client invokes an object that is located on a remote server.
- A remote object reference is passed in the invocation message to specify which object is to be invoked.
- Remote object references must be unique over space and time.

Remote Object References

- In general, may be many processes hosting remote objects, so remote object referencing must be unique among all of the processes in the various computers in a distributed system.
- generic format for remote object references is shown in **Figure 11**.




Figure 11. Representation of a remote object references

Remote Object References

- internet address/port number: process which created object
- time: creation time
- object number: local counter, incremented each time an object is created in the creating process
- interface: how to access the remote object (if object reference is passed from one client to another)

4.4 Client-Server Communication

- The client-server communication is designed to support the **roles** and **message exchanges** in typical client-server interactions.
- In the normal case, **request-reply** communication is **synchronous** because the **client** process **blocks** until the reply arrives from the server.
- **Asynchronous request-reply** communication is an alternative that is useful where **clients won't block** & can afford to retrieve replies later.

Client-Server Communication

- Protocol often built over UDP datagram's
 - UDP protocol avoids unnecessary **overheads** associated with **TCP**(Stream) protocol
 1. **acknowledgements** are **redundant**, since requests are followed by **replies**;(piggybacked Ack's)
 2. Avoidance of **connection establishment** overhead which involves two extra pairs of msg's.
 3. No need for **flow control** due to small amounts of data (arguments/results) are transferred
-

Client-Server Communication

- The request-reply protocol was based on a trio of communication primitives: **doOperation**, **getRequest**, and **sendReply** shown in Figure 12.
- Illustrates RMI example




Figure 12. Request-reply communication

Client-Server Communication

- **Figure 13** outlines the three communication primitives.

public byte[] doOperation (RemoteObjectRef o, int methodId, byte[] arguments)

sends a request message to the remote object and returns the reply.

The arguments specify the remote object, the method to be invoked and the arguments of that method.

public byte[] getRequest ();

acquires a client request via the server port.

public void sendReply (byte[] reply, InetAddress clientHost, int clientPort);

sends the reply message *reply* to the client at its Internet address and port.

Figure 13. Operations of the request-reply protocol

Client-Server Communication

public byte[] doOperation (RemoteObjectRef o, int methodId, byte[] arguments)

sends a request message to the remote object and returns the reply.

The arguments specify the remote object, the method to be invoked and the arguments of that method.

public byte[] getRequest ();

acquires a client request via the server port.

public void sendReply (byte[] reply, InetAddress clientHost, int clientPort);

sends the reply message *reply* to the client at its Internet address and port.

Figure 13. Operations of the request-reply protocol

Client-Server Communication

Three Primitives:

1.doOperation method

- used by the client to invoke remote operations
- Arguments → object & method to be invoked
- Its result is an RMI Reply
- Client calling doOperation marshals the arguments into an array of bytes & unmarshals the results from the array of bytes
- Client doOperation is blocked until remote object in the server performs the requested operation & transmits a reply msg back

2.GetRequest

- -used by a server process to acquire service requests
- -when server has invoked the method in the object it then uses

3.SendReply is used to send reply to client.

- -when reply msg is received doOperation is unblocked & client continues to execute

Client-Server Communication

- The **information** to be transmitted **in a request message** or a **reply message** is shown in Figure 14.

| | |
|-----------------|----------------------------------|
| messageType | <i>int (0=Request, 1= Reply)</i> |
| requestId | <i>int</i> |
| objectReference | <i>RemoteObjectRef</i> |
| methodId | <i>int or Method</i> |
| arguments | <i>// array of bytes</i> |

Figure 14. Request-reply message structure

Client-Server Communication

- In a protocol message
 - The **first field** indicates whether the message is a **request or a reply** message.
 - The **second field request id** contains a message identifier.
 - A **message identifier** consists of **two parts**:
 - ❖ A **requestId**, which is taken from an increasing sequence of integers by the sending process
 - ❖ An **identifier for the sender process**, for example its port and Internet address
 - The **third field** is a **remote object reference** .
 - The **forth field** is an **identifier for the method** to be invoked followed by **arguments**

Client-Server Communication

- Failure model of the request-reply protocol
 - If these **three primitives** are implemented over UDP they have the same **communication failures**

- Omission failure (link failures, drops/losses, missed/corrupt addresses)
- Out-of-order delivery
- Node/process down

Solved by

- Timeouts with retrans until reply is received/confirmed
- Discards of repeated requests by requestId (by server process)
- On lost reply messages, server repeats idempotent operations(eg.adding an element to set)
- Maintain history (reqid, message, client-id) or buffer replies and retrans – memory intensive

Client-Server Communication

- RPC exchange protocols(failure handling)
 - Three protocols are used for implementing various types of RPC.
 - ❖ The request (R) protocol.
 - ❖ The request-reply (RR) protocol.
 - ❖ The request-reply-acknowledge (RRA) protocol.

(Figure 15)

Client-Server Communication

| <i>Name</i> | <i>Messages sent by</i> | | |
|-------------|-------------------------|---------------|--------------------------|
| | <i>Client</i> | <i>Server</i> | <i>Client</i> |
| R | <i>Request</i> | | |
| RR | <i>Request</i> | <i>Reply</i> | |
| RRA | <i>Request</i> | <i>Reply</i> | <i>Acknowledge reply</i> |

Figure 15. RPC exchange protocols

Client-Server Communication

- In the **R protocol**, a single request message is sent by the client to the server.
- The R protocol may be used when there is no value to be returned from the remote method.
- The **RR protocol** is useful for most client-server exchanges because it is based on request-reply protocol.
- **RRA protocol** is based on the exchange of three messages: request-reply-acknowledge reply.

Client-Server Communication

- HTTP: an example of a request-reply protocol
 - HTTP is a request-reply protocol for the exchange of network resources between web clients and web servers.
 - Client requests specify a URL that includes DNS Host name+Port no.+resource identifier on that port
 - HTTP Allows→Content Negotiation & Authentication
 - Content negotiation→ negotiating for appropriate data representations between client & server
 - Password Style Authentication

Client-Server Communication

- HTTP protocol steps for C/S interaction:
 - ❖ Connection establishment between client and server at the default server port or at a port specified in the URL
 - ❖ client sends a request
 - ❖ server sends a reply
 - ❖ connection closure

Client-Server Communication

- Need to establish & close connection for every request-reply exchange is expensive
- Request & reply are marshalled into msgs as ASCII text
- Resources can have MIME(Multipurpose Internet Mail Extension)-like structures in arguments and results
- Data is prefixed with Mime type so that recipient will know how to handle it
- **Mime type** specifies a type and a subtype, for example:
 - ❖ text/plain, text/html, image/gif, image/jpeg

Client-Server Communication

- HTTP methods
- Client Rqst=method+URL

➤ GET

- ❖ Requests the resource, identified by URL as argument.
- ❖ If the URL refers to data, then the web server replies by returning the data
- ❖ If the URL refers to a program, then the web server runs the program and returns the output to the client.

| <i>method</i> | <i>URL</i> | <i>HTTP version</i> | <i>headers</i> | <i>message body</i> |
|---------------|--------------------------------|---------------------|----------------|---------------------|
| GET | //www.dcs.qmw.ac.uk/index.html | HTTP/ 1.1 | | |

Figure 16. HTTP request message

Client-Server Communication

➤ HEAD

- ❖ This method is similar to GET, but only meta data on resource is returned (like date of last modification, type, and size)
- ❖ i.e status line

Client-Server Communication

➤ POST

- ❖ Specifies the URL of a resource (for instance, a server program) that can deal with the data supplied with the request.
- ❖ This method is designed to deal with:
 - Providing a block of data to a data-handling process
 - Posting a message to a bulletin board, mailing list or news group.
 - Extending a dataset with an append operation

Client-Server Communication

➤ PUT

- ❖ Supplied data to be stored in the given URL as its identifier.

➤ DELETE

- ❖ The server deletes an identified resource by the given URL on the server.

➤ TRACE

- ❖ The server sends back the request message

➤ OPTIONS

- ❖ A server supplies the client with a list of methods.
- ❖ It allows to be applied to the given URL

Client-Server Communication

- A reply message specifies
 - ❖ The protocol version
 - ❖ A status code
 - ❖ Reason
 - ❖ Some headers
 - ❖ An optional message body

| <i>HTTP version</i> | <i>status code</i> | <i>reason</i> | <i>headers</i> | <i>message body</i> |
|---------------------|--------------------|---------------|----------------|---------------------|
| HTTP/1.1 | 200 | OK | | resource data |

Figure 17. HTTP reply message

Client-Server Communication

■ Status codes

- 100 block → Informational → Eg. 103 – checkpoint
- 200 block → Success → Eg. 200-OK, 201-created
- 300 block → Redirection → Eg. 302-Found, 304-Not Modified
- 400 block → Client Error → Eg. 404-Not Found 408-Request Timeout
- 500 block → Server error → Eg. 500-Internal Server Error, 502-Bad Gateway, 503-Service Unavailable