

1.5 A Simple Daytime Server

We can write a simple version of a TCP daytime server, which will work with the client from [Section 1.2](#). We use the wrapper functions that we described in the previous section and show this server in [Figure 1.9](#).

Figure 1.9 TCP daytime server.

intro/daytimetcpsrv.c

```

1 #include      "unp.h".
2 #include      <time.h>

3 int
4 main(int argc, char **argv)
5 {
6     int      listenfd, connfd;
7     struct sockaddr_in servaddr;
8     char      buff[MAXLINE];
9     time_t ticks;

10    listenfd = Socket(AF_INET, SOCK_STREAM, 0);

11    bzeros(&servaddr, sizeof(servaddr));
12    servaddr.sin_family = AF_INET;
13    servaddr.sin_addr.s_addr = htonl(INADDR_ANY);
14    servaddr.sin_port = htons(13); /* daytime server */

15    Bind(listenfd, (SA *) &servaddr, sizeof(servaddr));

16    Listen(listenfd, LISTENQ);

17    for ( ; ; ) {
18        connfd = Accept(listenfd, (SA *) NULL, NULL);

19        ticks = time(NULL);
20        snprintf(buff, sizeof(buff), "%.24s\r\n", ctime(&ticks));
21        Write(connfd, buff, strlen(buff));

22        Close(connfd);
23    }
24 }
```

Create a TCP socket

10 The creation of the TCP socket is identical to the client code.

Bind server's well-known port to socket

11–15 The server's well-known port (13 for the daytime service) is bound to the socket by filling in an Internet socket address structure and calling `bind`. We specify the IP address as `INADDR_ANY`, which allows the server to accept a client connection on any interface, in case the server host has multiple interfaces. Later we will see how we can restrict the server to accepting a client connection on just a single interface.

Convert socket to listening socket

16 By calling `listen`, the socket is converted into a listening socket, on which incoming connections from clients will be accepted by the kernel. These three steps, `socket`, `bind`, and `listen`, are the normal steps for any TCP server to prepare what we call the *listening descriptor* (`listenfd` in this example).

The constant `LISTENQ` is from our `unp.h` header. It specifies the maximum number of client connections that the kernel will queue for this listening descriptor. We say much more about this queueing in [Section 4.5](#).

Accept client connection, send reply

17–21 Normally, the server process is put to sleep in the call to `accept`, waiting for a client connection to arrive and be accepted. A TCP

connection uses what is called a *three-way handshake* to establish a connection. When this handshake completes, `accept` returns, and the return value from the function is a new descriptor (`connfd`) that is called the *connected descriptor*. This new descriptor is used for communication with the new client. A new descriptor is returned by `accept` for each client that connects to our server.

The style used throughout the book for an infinite loop is

```
for ( ; ; ) {  
    . . .  
}
```

The current time and date are returned by the library function `time`, which returns the number of seconds since the Unix Epoch: 00:00:00 January 1, 1970, Coordinated Universal Time (UTC). The next library function, `ctime`, converts this integer value into a human-readable string such as

```
Mon May 26 20:58:40 2003
```

A carriage return and linefeed are appended to the string by `snprintf`, and the result is written to the client by `write`.

If you're not already in the habit of using `snprintf` instead of the older `sprintf`, now's the time to learn. Calls to `sprintf` cannot check for overflow of the destination buffer. `snprintf`, on the other hand, requires that the second argument be the size of the destination buffer, and this buffer will not overflow.

`snprintf` was a relatively late addition to the ANSI C standard, introduced in the version referred to as *ISO C99*. Virtually all vendors provide it as part of the standard C library, and many freely available versions are also available. We use `snprintf` throughout the text, and we recommend using it instead of `sprintf` in all your programs for reliability.

It is remarkable how many network break-ins have occurred by a hacker sending data to cause a server's call to `sprintf` to overflow its buffer. Other functions that we should be careful with are `gets`, `strcpy`, and `strncpy`, normally calling `fgets`, `strncat`, and `strncpy` instead. Even better are the more recently available functions `strlcat` and `strlcpy`, which ensure the result is a properly terminated string. Additional tips on writing secure network programs are found in Chapter 23 of [Garfinkel, Schwartz, and Spafford 2003].

Terminate connection

²² The server closes its connection with the client by calling `close`. This initiates the normal TCP connection termination sequence: a FIN is sent in each direction and each FIN is acknowledged by the other end. We will say much more about TCP's three-way handshake and the four TCP packets used to terminate a TCP connection in [Section 2.6](#).

As with the client in the previous section, we have only examined this server briefly, saving all the details for later in the book. Note the following points:

- As with the client, the server is protocol-dependent on IPv4. We will show a protocol-independent version that uses the `getaddrinfo` function in [Figure 11.13](#).
- Our server handles only one client at a time. If multiple client connections arrive at about the same time, the kernel queues them, up to some limit, and returns them to `accept` one at a time. This daytime server, which requires calling two library functions, `time` and `ctime`, is quite fast. But if the server took more time to service each client (say a few seconds or a minute), we would need some way to overlap the service of one client with another client.
- The server that we show in [Figure 1.9](#) is called an *iterative server* because it iterates through each client, one at a time. There are numerous techniques for writing a *concurrent server*, one that handles multiple clients at the same time. The simplest technique for a concurrent server is to call the Unix `fork` function ([Section 4.7](#)), creating one child process for each client. Other techniques are to use threads instead of `fork` ([Section 26.4](#)), or to pre-fork a fixed number of children when the server starts ([Section 30.6](#)).
- If we start a server like this from a shell command line, we might want the server to run for a long time, since servers often run for as long as the system is up. This requires that we add code to the server to run correctly as a Unix *daemon*: a process that can run in the background, unattached to a terminal. We will cover this in [Section 13.4](#).

[[Team LiB](#)]

◀ PREVIOUS

NEXT ▶